

Chapter 3: Processes



Chapter 3: Processes

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ Inter-Process Communication (IPC)
- ❑ IPC in Shared-Memory Systems
- ❑ IPC in Message-Passing Systems
- ❑ Examples of IPC Systems
- ❑ Communication in Client-Server Systems



Objectives

- ❑ Identify the separate **components of a process** and illustrate how they are represented and scheduled in an operating system.
- ❑ Describe **how processes are created and terminated** in an operating system, including developing programs using the appropriate system calls that perform these operations.
- ❑ Describe and contrast **inter-process communication** using shared memory and message passing.
- ❑ Design *programs that uses pipes and POSIX shared memory* to perform inter-process communication.
- ❑ Describe **client-server communication** using sockets and remote procedure calls.
- ❑ Design *kernel modules* that interact with the Linux operating system.



Process Concept

- ❑ An operating system executes a variety of programs that run as processes
- ❑ **Process** – a program in execution; process execution must progress in sequential fashion
- ❑ Multiple parts
 - The *program code*, also called *text section*
 - Current activity including *program counter*, and *processor registers*
 - *Stack section* containing temporary data
 - ▶ Function parameters, return addresses, local variables
 - *Data section* containing global variables
 - *Heap section* containing memory dynamically allocated during run time




Process Concept (Cont.)

- ❑ *Program* is *passive* entity stored on disk (e.g., *executable file*)
- ❑ *Process* is *active* entity
 - Program becomes process when executable file loaded into memory
- ❑ *Execution of program* can be started via GUI mouse clicks, command line (CLI) entry of its name, etc.
- ❑ One program can be several processes
 - E.g., Consider multiple users executing the same program

`#ps -aux`



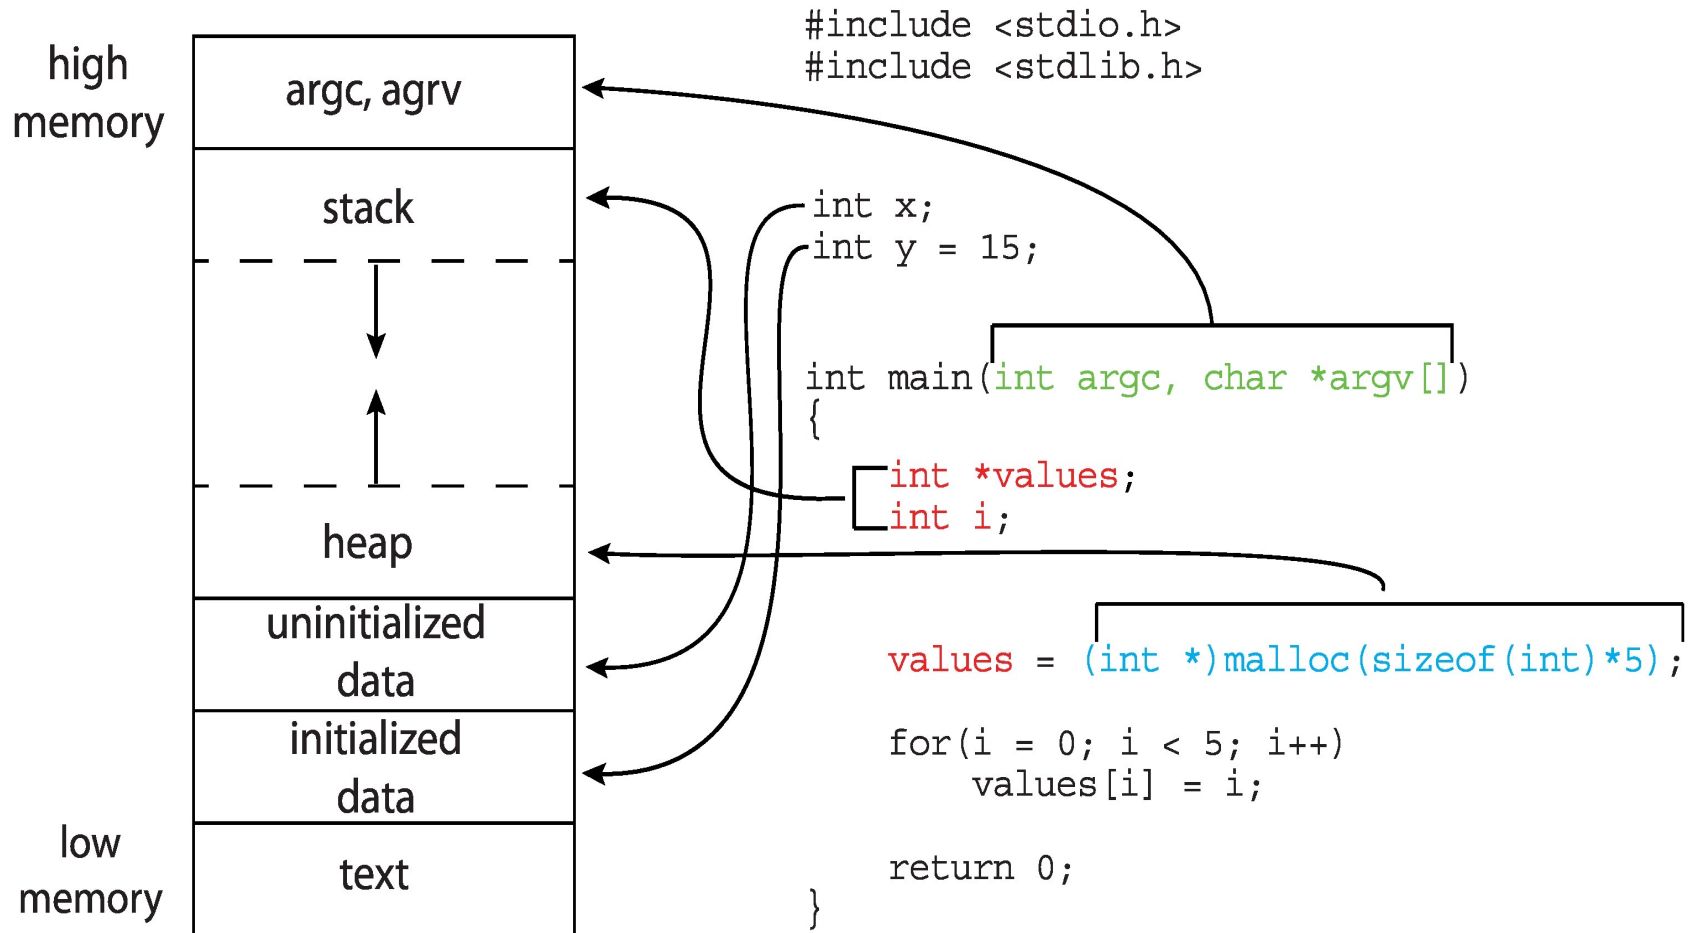
Process in Memory



#size <pid>



Memory Layout of a C Program

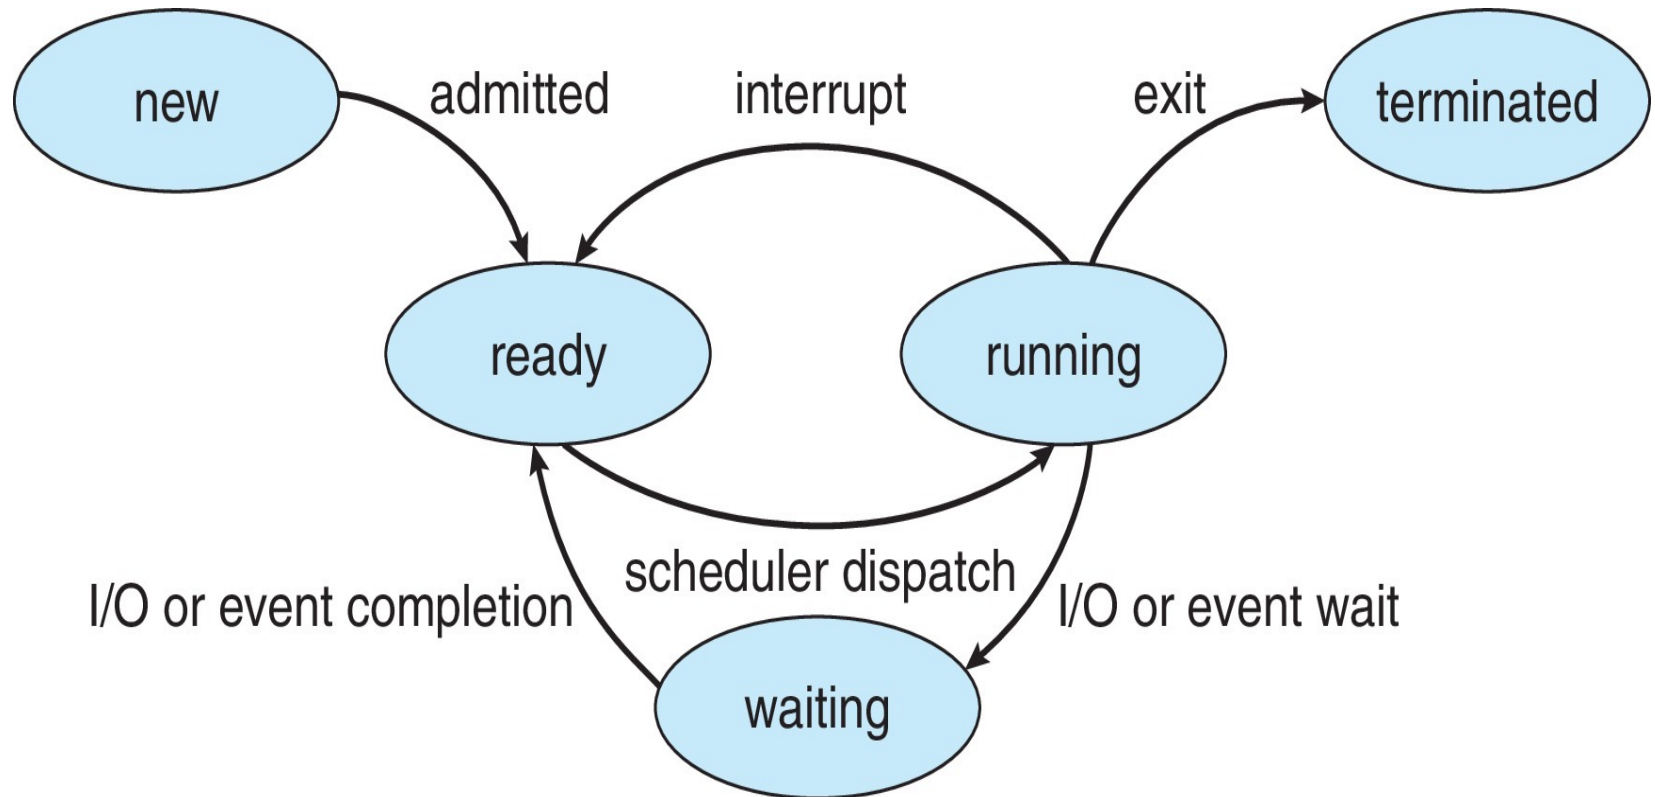


Process State

- ❑ As a process executes, it changes *state*
 - *New* – The process is being created
 - *Running* – Instructions are being executed
 - *Waiting* – The process is waiting for some event to occur
 - *Ready* – The process is waiting to be assigned to a processor
 - *Terminated* – The process has finished execution



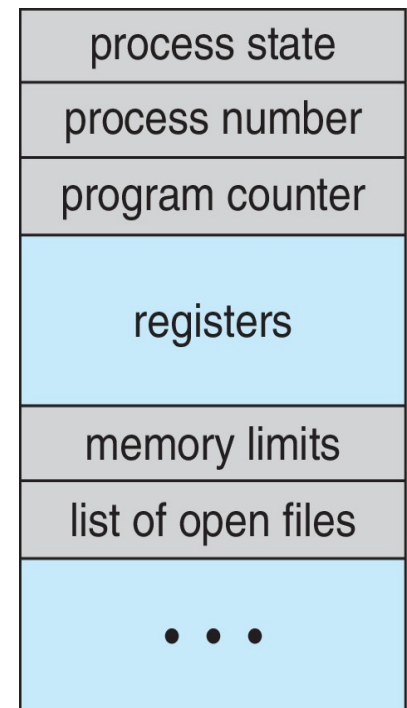
Diagram of Process State



Process Control Block (PCB)

❑ **Process Control Block (PCB)** – Information associated with each process, also called **Task Control Block (TCB)**, includes:

- *Process state* – running, waiting, etc.
- *Process number* – identity of the process
- *Program counter* – location of instruction to next execute
- *CPU registers* – contents of all process-centric registers
- *CPU scheduling info* – priorities, scheduling queue pointers
- *Memory-management information* – memory allocated to the process
- *Accounting information* – CPU used, clock time elapsed since start, time limits
- *I/O status information* – I/O devices allocated to process, list of open files



Threads

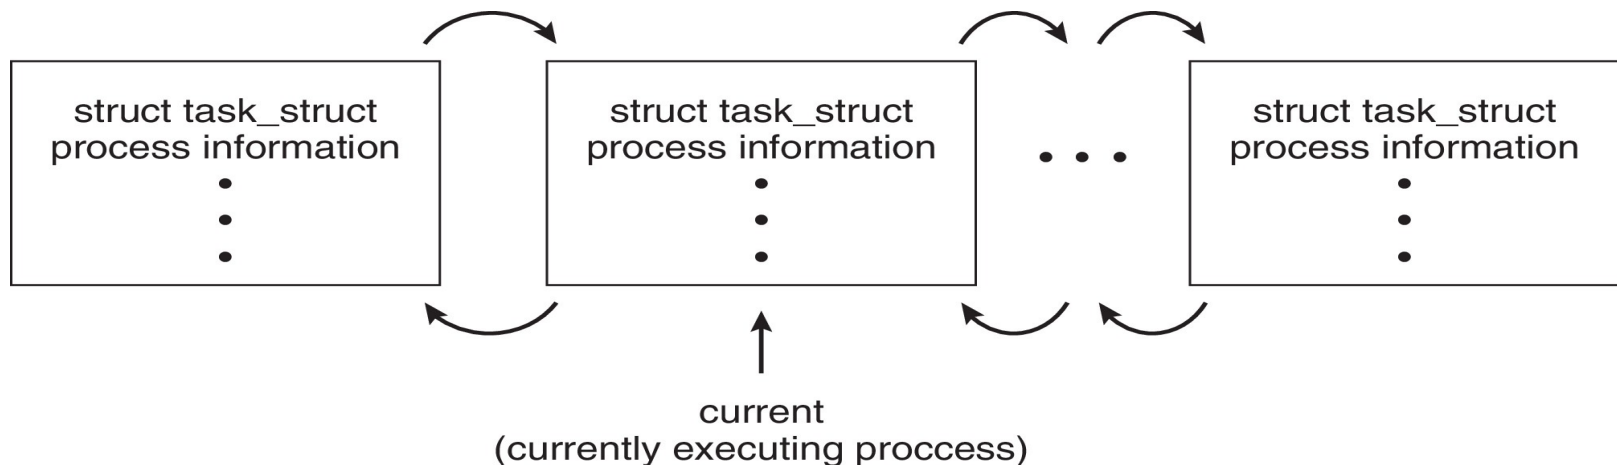
- ❑ So far, process has a *single thread* of execution
- ❑ Consider having *multiple program counters per process*
 - Multiple locations can execute at once
 - ▶ Multiple threads of control → *threads*
- ❑ Must then have *storage for thread* details
- ❑ Multiple program counters in PCB



Process Representation in Linux

❑ Represented by the C structure `task_struct`

```
pid t_pid;                /* process identifier */
long state;               /* state of the process */
unsigned int time_slice;  /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm;      /* address space of this process */
```

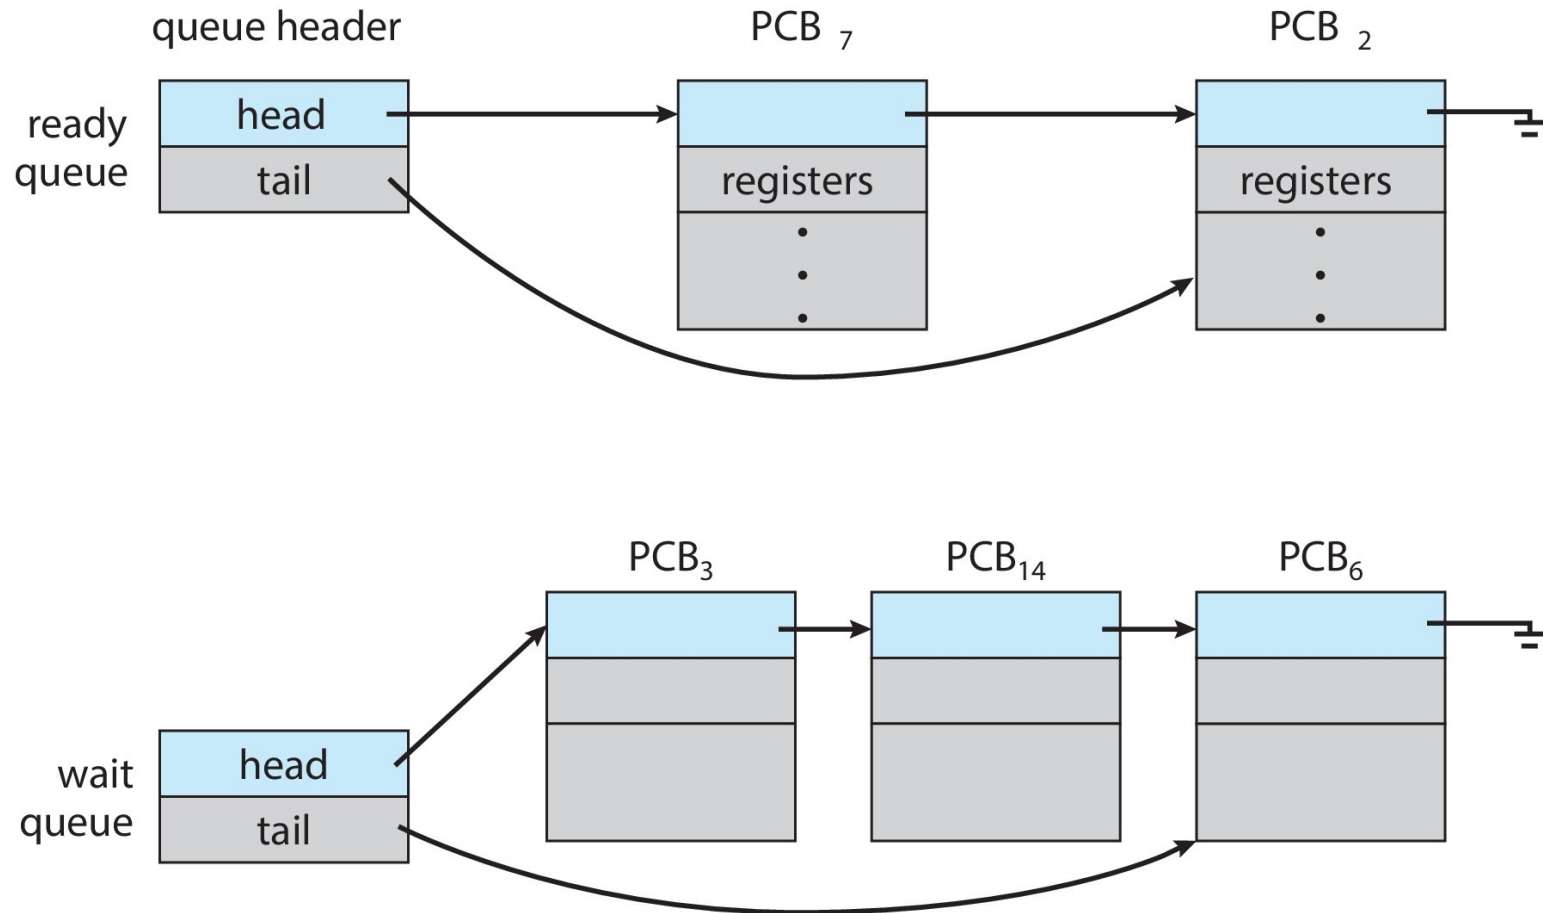


Process Scheduling


- ❑ Maximize CPU use → quickly switch processes onto CPU core
- ❑ *Process scheduler* selects one process among available (ready) processes for next execution on CPU core
- ❑ Maintains *scheduling queues* of processes
 - *Ready queue* – set of all processes residing in main memory, ready and waiting to execute
 - *Wait queues* – set of processes waiting for an event (e.g., I/O)
- ❑ Processes migrate among the various queues



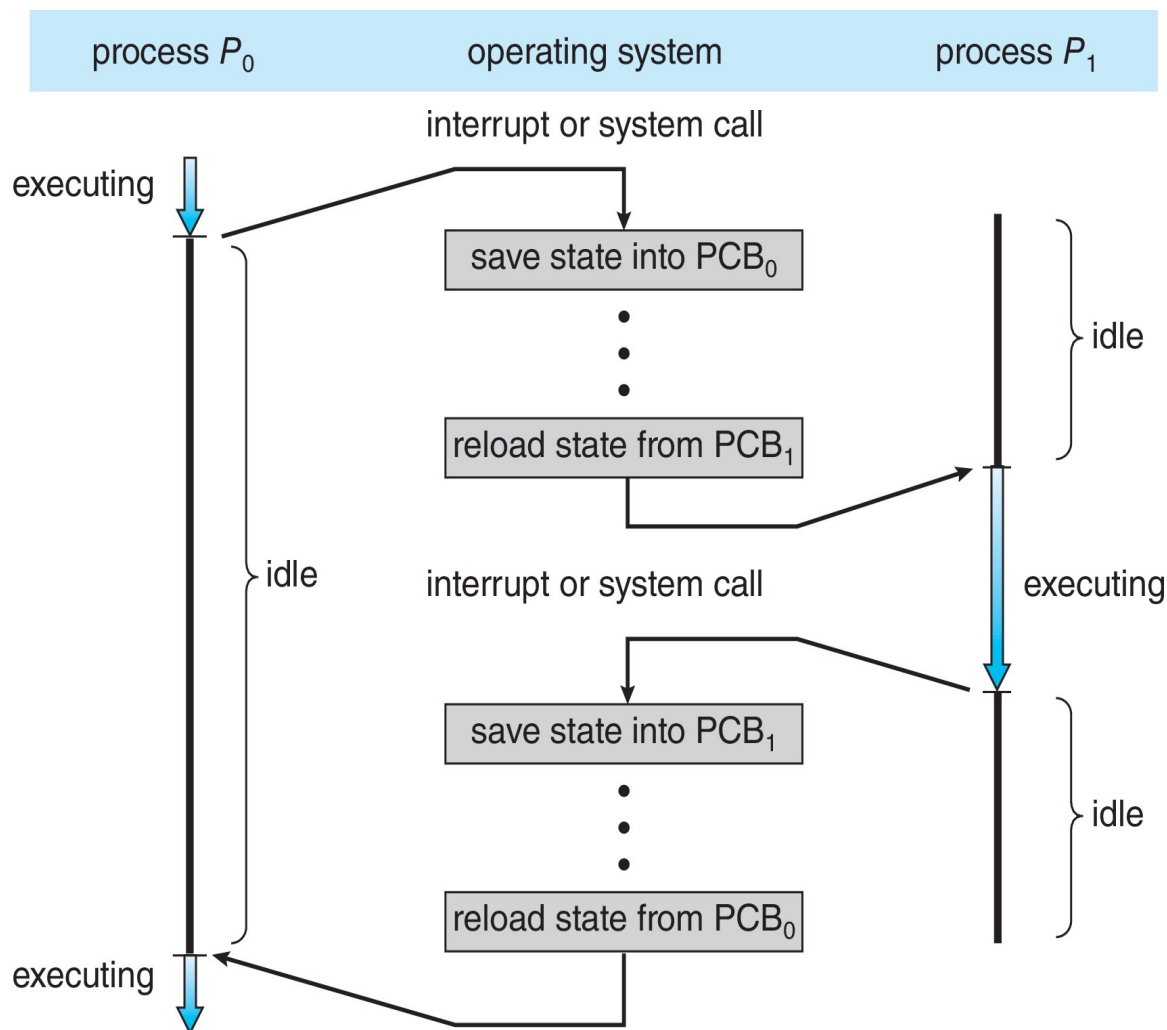
Ready and Wait Queues



Representation of Process Scheduling



CPU Switch from Process to Process



- A *context switch* occurs when the CPU switches from one process to another.

Context Switch

- ❑ When CPU switches to another process, the system must *save the state* of the old process and load the *saved state* for the new process via a *context switch*
- ❑ *Context* of a process represented in the **PCB**
- ❑ Context-switch time is *overhead*, the system does no useful work while switching
 - The more complex the OS and the PCB, the longer the context switch
- ❑ *Time* dependent on hardware support
 - Some hardware provides *multiple sets of registers per CPU*, multiple contexts loaded at once



Multitasking in Mobile Systems

- ❑ Some *mobile systems* (e.g., early version of iOS) allow only one process to run, others suspended
- ❑ Due to screen real state, user interface limits **iOS** provides for a
 - Single *foreground process* – controlled via user interface
 - Multiple *background processes* – in memory, running, but not on the display, and with limits
 - *Limits* include single, short task, receiving notification of events, specific long-running tasks like audio playback
- ❑ **Android** runs foreground and background, with fewer limits
 - Background process uses a *service* to perform tasks
 - Service can keep running even if background process is suspended
 - Service has no user interface, small memory use



Operations on Processes

- ❑ System must provide mechanisms for:
 - process creation
 - process termination



Process Creation

- ❑ *Parent processes* create *children processes*, which, in turn create other processes, forming a *tree of processes*
- ❑ Process identified and managed via a **Process Identifier (PID)**
- ❑ **Resource sharing options**
 - Parent and children share *all* resources
 - Children share *subset* of parent's resources
 - Parent and child share *no* resources



Process Creation (Cont.)

❑ Execution options


- Parent and children execute concurrently
- Parent waits until children terminate

❑ Address space

- Child duplicate of parent
- Child has a program loaded into it



A Tree of Processes in Linux



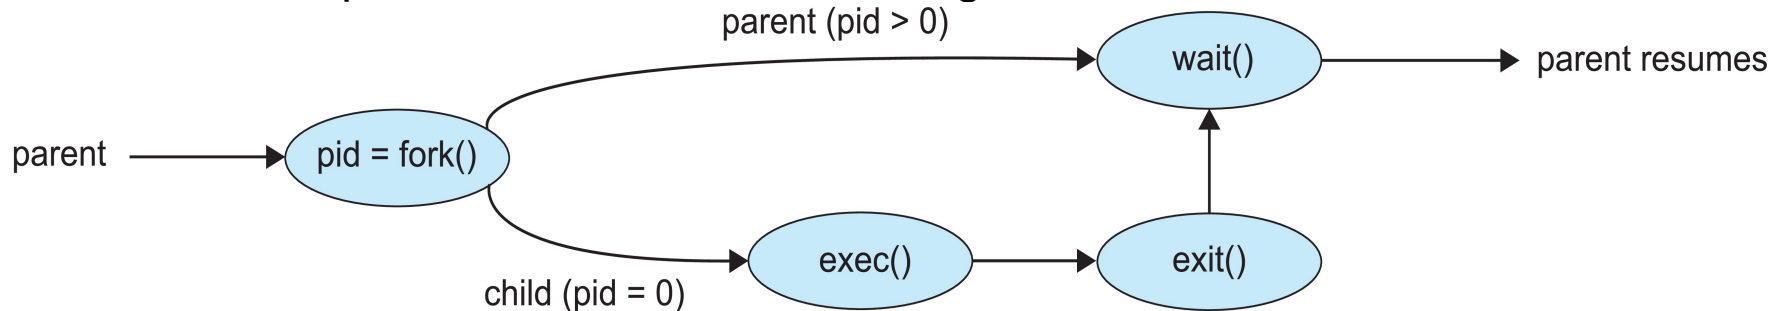
#pstree



Process Creation (Cont.)

❑ UNIX examples

- **system()** to execute a command from within a program
- **fork()** system call creates new process
- **exec()** system call used after a **fork()** to replace the process' memory space with a new program
- Parent process calls **wait()** waiting for the child to terminate



system() system call

```
#include <stdlib.h>
int system(const char *command);
```

- ❑ hands the argument command to the command interpreter sh: `$/bin/sh -c <command>`
 - Ex: `$/bin/sh -c ls`
- ❑ If command is **NULL**, system() return a nonzero value if a shell is available, or 0 if no shell is available

```
#include <stdlib.h>
```

```
int main ( )
{
    int return_value ;
    return_value = system ( "ls ." );
    return return_value;
}
```

could not be created or its status could not be

```
Vans-MacBook-Air:c-examp ltvan$ ./systemsimple
2.2.c.rtf      ex          fork2.c      pid
addrspace     ex.c        forkp        pid.cpp
```



C Program Forking A Separate Process: fork()

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```



Example

❑

How many processes will be generated?

❑ int main(){

❑ printf("Hello \n");

❑ 1. fork();

❑ 2. fork();

❑ printf("Hello \n");

❑ return 0;

❑ }



Process Termination

- ❑ Process executes *last statement* and then asks the operating system to delete it using the `exit()` system call.
 - Returns status data from child to parent (via `wait()`)
 - Process' resources are deallocated by operating system
- ❑ Parent may terminate the execution of children processes using the `abort()` system call. Some reasons for doing so:
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - The parent is exiting and the *operating systems does not allow a child to continue if its parent terminates*



Process Termination (Cont.)

- ❑ Some operating systems do not allow child to exist if its parent has terminated. *If a process terminates, then all its children must also be terminated.*
 - **Cascading termination:** All children, grandchildren, etc. are terminated
 - The termination is initiated by the operating system
- ❑ The parent process may wait for termination of a child process by using the `wait()` system call. The call returns status information and the **pid** of the terminated process
 - ✓ `pid = wait(&status);`
 - ✓ `waitpid()` suspends execution of the calling process until a child specified by *pid* argument has changed state
- ❑ If no parent waiting (did not invoke `wait()`), process is a *zombie*
- ❑ If parent terminated without invoking `wait()`, process is an *orphan*



waitpid() example

```
#include <sys/types.h>
#include <sys/wait.h>
```

```
pid_t waitpid(pid_t pid, int
*status, int options);
```

- `pid < -1`: wait for any child process whose process group ID is equal to the absolute value of `pid`.
- `pid = 1`: wait for any child process
- `pid = 0`: wait for any child process whose process group ID is equal to that of the calling process
- `pid > 0`: wait for the child whose process ID is equal to the value of `pid`

```
#include<stdio.h>
#include<sys/wait.h>
#include<unistd.h>

int main(){
    pid_t pid=fork();
    if(pid==0){
        printf("Child proc id = %d, groupid = %d \n",
getpid(),getpgrp());
    }
    else if(pid>0){
        waitpid(pid,NULL,0);
        printf("Parent proc id = %d \n",getpid());
    }
    return 0;
}
```

`wait(&status)` is equivalent to
`waitpid(-1, &status, 0);`

```
Vans-MacBook-Air:c-examp ltvan$ ./forkwpidr
Child proc id = 9716, groupid = 9715
Parent proc id = 9715
```

wait(&status) example

```
#include <sys/types.h>
#include <sys/wait.h>
pid_t wait(int *status);
```

```
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>
int main (int argc, char *argv[]){
    int    return_code;
    int status;
    /* create a new process */
    return_code = fork();
    if (return_code > 0){
```

```
        int id = wait(&status);
        printf("Parent process: child's pid = %d with returned value = %d \n", id, WEXITSTATUS(status));
        printf("child's pid = %d with raw returned value = %d \n", id, status);
        printf("Pid of child %d \n", return_code);
        return 0;
```

```
    }
    else if (return_code == 0){
        printf("This is child process \n");
        return 10; //exit(10);
    }
    else {
        printf("Fork error\n");
        return 1;
    }
}
```



Exercise

- ✓ Given the following code:

```
int main(){  
    int i;  
    for(i=0; i<n;i++)  
        fork();  
    return 0;  
}
```

How many processes will be generated if $n = 2$?




Importance Hierarchy of Android Process

- ❑ *Mobile operating systems* often have to terminate processes to reclaim system resources such as memory. From most to least important:
 - ▲ Foreground process
 - ▲ Visible process
 - ▲ Service process
 - ▲ Background process
 - ▲ Empty process
- ❑ Android will begin terminating processes that are least important.



Multiprocess Architecture – Chrome Browser

- ❑ Many web browsers ran as a single process (some still do)
 - If one web site causes trouble, entire browser can hang or crash
- ❑ Google Chrome Browser is multiprocess with 3 different types of processes:
 - *Browser process* manages user interface, disk and network I/O
 - *Renderer process* renders web pages, deals with HTML, JavaScript. A new renderer created for each website opened
 - ▶ Runs in *sandbox* restricting disk and network I/O, minimizing effect of security exploits
 - *Plug-in process* for each type of plug-in



Each tab represents a separate process.



Inter-Process Communication (IPC)


- ❑ Processes within a system may be *independent* or *cooperating*
 - *Independent process* does not share data with any other processes executing in the system
 - *Cooperating process* can affect or be affected by other processes, including sharing data
- ❑ Reasons for cooperating processes:
 - Information sharing
 - Computation speed-up
 - Modularity
 - Convenience
- ❑ Cooperating processes need **Inter-Process communication (IPC)**




Communication Models

Two models of IPC

- *Shared memory*
- *Message passing*



(a)



(b)

Inter-Process Communication – Shared Memory

- ❑ An *area of memory shared among the processes* that wish to communicate
- ❑ The communication is *under the control of the users processes*, not the operating system.
- ❑ Major issues is to provide mechanism that will allow the user processes to *synchronize their actions* when they access shared memory.



Producer-Consumer Problem

- ❑ *Producer-Consumer relationship*
- ❑ Paradigm for cooperating processes, *producer process* produces information that is consumed by a *consumer process*
 - *unbounded-buffer* places no practical limit on the size of the buffer
 - *bounded-buffer* assumes that there is a fixed buffer size



Bounded-Buffer – Shared-Memory Solution

❑ Shared data

```
#define BUFFER_SIZE 10

typedef struct {

    . . .

} item;

item buffer[BUFFER_SIZE];

int in = 0;

int out = 0;
```

❑ Solution is correct, but can only use **BUFFER_SIZE-1** elements



Producer Process – Shared Memory

```
item next_produced;

while (true) {

    /* produce an item in next produced */
    while (((in + 1) % BUFFER_SIZE) == out)

        ; /* do nothing */

    buffer[in] = next_produced;

    in = (in + 1) % BUFFER_SIZE;

}
```



Consumer Process – Shared Memory

```
item next_consumed;

while (true) {
    while (in == out)

        ; /* do nothing */
    next_consumed = buffer[out];

    out = (out + 1) % BUFFER_SIZE;

    /* consume the item in next consumed */
}
```



Inter-Process Communication – Message Passing

- ❑ Mechanism for processes to communicate and to synchronize their actions
- ❑ Message system – processes communicate with each other without resorting to shared variables
- ❑ IPC facility provides two operations:
 - send(message)
 - receive(message)
- ❑ The message size is either fixed or variable



Message Passing (Cont.)

- ❑ If processes P and Q wish to communicate, they need to:
 - Establish a communication link between them
 - Exchange messages via *send/receive*
- ❑ Implementation issues:
 - How are links established?
 - Can a link be associated with more than two processes?
 - How many links can there be between every pair of communicating processes?
 - What is the capacity of a link?
 - Is the size of a message that the link can accommodate fixed or variable?
 - Is a link unidirectional or bi-directional?



Direct Communication

- ❑ Processes must name each other explicitly:
 - `send(P, message)` – send a message to process P
 - `receive(Q, message)` – receive a message from process Q
- ❑ Properties of communication link
 - Links are established automatically
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there exists exactly one link
 - The link may be unidirectional, but is usually bi-directional



Indirect Communication

- ❑ Messages are directed and received from *mailboxes* (also referred to as *ports*)
 - Each mailbox has a *unique ID*
 - Processes can communicate *only if they share a mailbox*
- ❑ Properties of communication link
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bi-directional



Indirect Communication (Cont.)

❑ Operations

- create a new mailbox (or port)
- send and receive messages through mailbox
- destroy a mailbox

❑ **Primitives** are defined as:

- **send**(*A*, *message*) – send a message to mailbox A
- **receive**(*A*, *message*) – receive a message from mailbox A



Indirect Communication (Cont.)

❑ Mailbox sharing

○ Example

- ▶ P_1 , P_2 , and P_3 share mailbox A,
- ▶ P_1 sends; P_2 and P_3 receive.
- ▶ Who gets the message?

○ Solutions

- ▶ Allow a link to be associated with at most two processes
- ▶ Allow only one process at a time to execute a receive operation
- ▶ Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was



Message Passing – Synchronization

- ❑ Message passing may be either *blocking* or *non-blocking*
- ❑ *Blocking* is considered *synchronous*
 - *Blocking send* – the sender is blocked until the message is received
 - *Blocking receive* – the receiver is blocked until a message is available
- ❑ *Non-blocking* is considered *asynchronous*
 - *Non-blocking send* – the sender sends the message and continue
 - *Non-blocking receive* – the receiver receives:
 - ▶ A valid message, or Null message
- ❑ Different combinations possible
 - If both send and receive are blocking, we have a *rendezvous*



Producer – Message Passing

```
message next_produced;  
  
while (true) {  
    /* produce an item in next_produced */  
  
    send(next_produced);  
}
```



Consumer – Message Passing

```
message next_consumed;  
  
while (true) {  
    receive(next_consumed)  
  
    /* consume the item in next_consumed */  
}
```



Buffering

- ❑ Queue of messages attached to the link.
- ❑ Implemented in one of three ways
 - *Zero capacity* – no messages are queued on a link
 - ▶ Sender must wait for receiver (rendezvous)
 - *Bounded capacity* – finite length of n messages
 - ▶ Sender must wait if link full
 - *Unbounded capacity* – infinite length
 - ▶ Sender never waits



Examples of IPC Systems - POSIX

❑ POSIX Shared Memory

- Process first creates shared memory segment

```
shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);
```

- Also used to open an existing segment
- Set the size of the object

```
ftruncate(shm_fd, 4096);
```

- Use `mmap()` to memory-map a file pointer to the shared memory object
- Reading and writing to shared memory is done by using the pointer returned by `mmap()`.


Pipes

- ❑ Acts as a conduit allowing two processes to communicate
- ❑ Issues:
 - Is communication unidirectional or bidirectional?
 - In the case of two-way communication, is it half or full-duplex?
 - Must there exist a relationship (e.g., *parent-child*) between the communicating processes?
 - Can the pipes be used over a network?
- ❑ *Ordinary pipes* – cannot be accessed from outside the process that created it. Typically, a parent process creates a pipe and uses it to communicate with a child process that it created.
- ❑ *Named pipes* – can be accessed without a parent-child relationship.



Ordinary Pipes

- ❑ *Ordinary Pipes* allow communication in standard producer-consumer style
 - *Producer* writes to one end (the *write-end* of the pipe)
 - *Consumer* reads from the other end (the *read-end* of the pipe)
- ❑ Ordinary pipes are therefore unidirectional
- ❑ Require *parent-child relationship* between communicating processes



Named Pipes

- ❑ *Named pipes* are more powerful than ordinary pipes
- ❑ Communication is bidirectional
- ❑ *No parent-child relationship* is necessary between the communicating processes
- ❑ Several processes can use the named pipe for communication
- ❑ Provided on both **UNIX** and **Windows** systems



Communications in Client-Server Systems


❑ Sockets

- A *socket* is defined as an endpoint for communication
- It is a concatenation of *IP address* and *port* – a number included at start of message packet to differentiate network services on a host
 - ▶ E.g., The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**
- Communication consists between a pair of sockets
- All ports below 1024 are *well known*, used for standard services
- Special IP address **127.0.0.1** (*loopback*) to refer to system on which process is running

❑ Remote Procedure Calls (RPC)



Socket Communication



Sockets in Java – Server

```
import java.net.*;
import java.io.*;

public class DateServer
{
    public static void main(String[] args) {
        try {
            ServerSocket sock = new ServerSocket(6013);

            /* now listen for connections */
            while (true) {
                Socket client = sock.accept();

                PrintWriter pout = new
                    PrintWriter(client.getOutputStream(), true);

                /* write the Date to the socket */
                pout.println(new java.util.Date().toString());

                /* close the socket and resume */
                /* listening for connections */
                client.close();
            }
        } catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```

❑ Three types of sockets

- *Connection-oriented (TCP)*
- *Connectionless (UDP)*
- **MulticastSocket** class– data can be sent to multiple recipients

❑ Consider this “Date” *server* in Java:



Sockets in Java – Client

```
import java.net.*;
import java.io.*;

public class DateClient
{
    public static void main(String[] args) {
        try {
            /* make connection to server socket */
            Socket sock = new Socket("127.0.0.1",6013);

            InputStream in = sock.getInputStream();
            BufferedReader bin = new
                BufferedReader(new InputStreamReader(in));

            /* read the date from the socket */
            String line;
            while ( (line = bin.readLine()) != null)
                System.out.println(line);

            /* close the socket connection*/
            sock.close();
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```

- ❑ The equivalent
“Date” *client*



Remote Procedure Calls

- ❑ **Remote Procedure Call (RPC)** abstracts procedure calls between processes on *networked systems*
 - Again uses *ports* for service differentiation
- ❑ **Stubs** – proxies for the actual procedure on the server and client sides
 - The *client-side stub* locates the server and *marshals* the parameters
 - The *server-side stub* receives this message, unpacks the marshalled parameters, and performs the procedure on the server
- ❑ On **Windows**, stub code compile from specification written in **Microsoft Interface Definition Language (MIDL)**



Remote Procedure Calls (Cont.)

- ❑ Data representation handled via **External Data Representation (XDR)** format to account for different architectures
 - E.g., *Big-endian* (Motorola) and *little-endian* (Intel x86)
- ❑ Remote communication has *more failure scenarios* than local
 - Messages can be delivered *exactly once* rather than *at most once*
- ❑ OS typically provides a *rendezvous* (or *matchmaker*) service to connect client and server



End of Chapter 3

